How the Internet Works

Network

What is the Internet?

- A global network of billions of interconnected devices that request and receive information
- Every device can be uniquely identified by an **IP (Internet Protocol)** address, just like you have your physical home address
- Many **network protocols** serve different purposes; think of them as "**network languages**"
- Within the scope of your home network, devices are assigned *local* IP addresses, and they are all "hidden" from the Internet with a single *external* IP address that is assigned to your main home router by the **ISP (Internet Service Provider)**

- ***So, what happens when you want to get to a website?***

DNS – the Domain Name System

- Let's say you enter a **domain name** (a human-readable address): `chatgpt.com`
- Devices can only navigate the Internet using IP addresses, which are numeric, like `192.168.0.2` for your home network or `172.64.155.209` for ChatGPT
- But how would your device know that `172.64.155.209` belongs to ChatGPT based on a single `chatgpt.com` name?
- This is where a **DNS** protocol comes in handy (port number **53**, and you'll see more on ports later) that helps look up an IP, just like you would look up a **phone in a phonebook**
- Your browser will first ask a DNS server (the network you join provides the IP address of that server automatically) to *resolve* the domain name and provide the IP address

DNS Record Type	Hostname	Value
A	chatgpt.com	172.64.155.209

DNS: A Quick Glance

Name server: think of it as the DNS server itself
TLD: top-level domain like `.com` or `.edu`

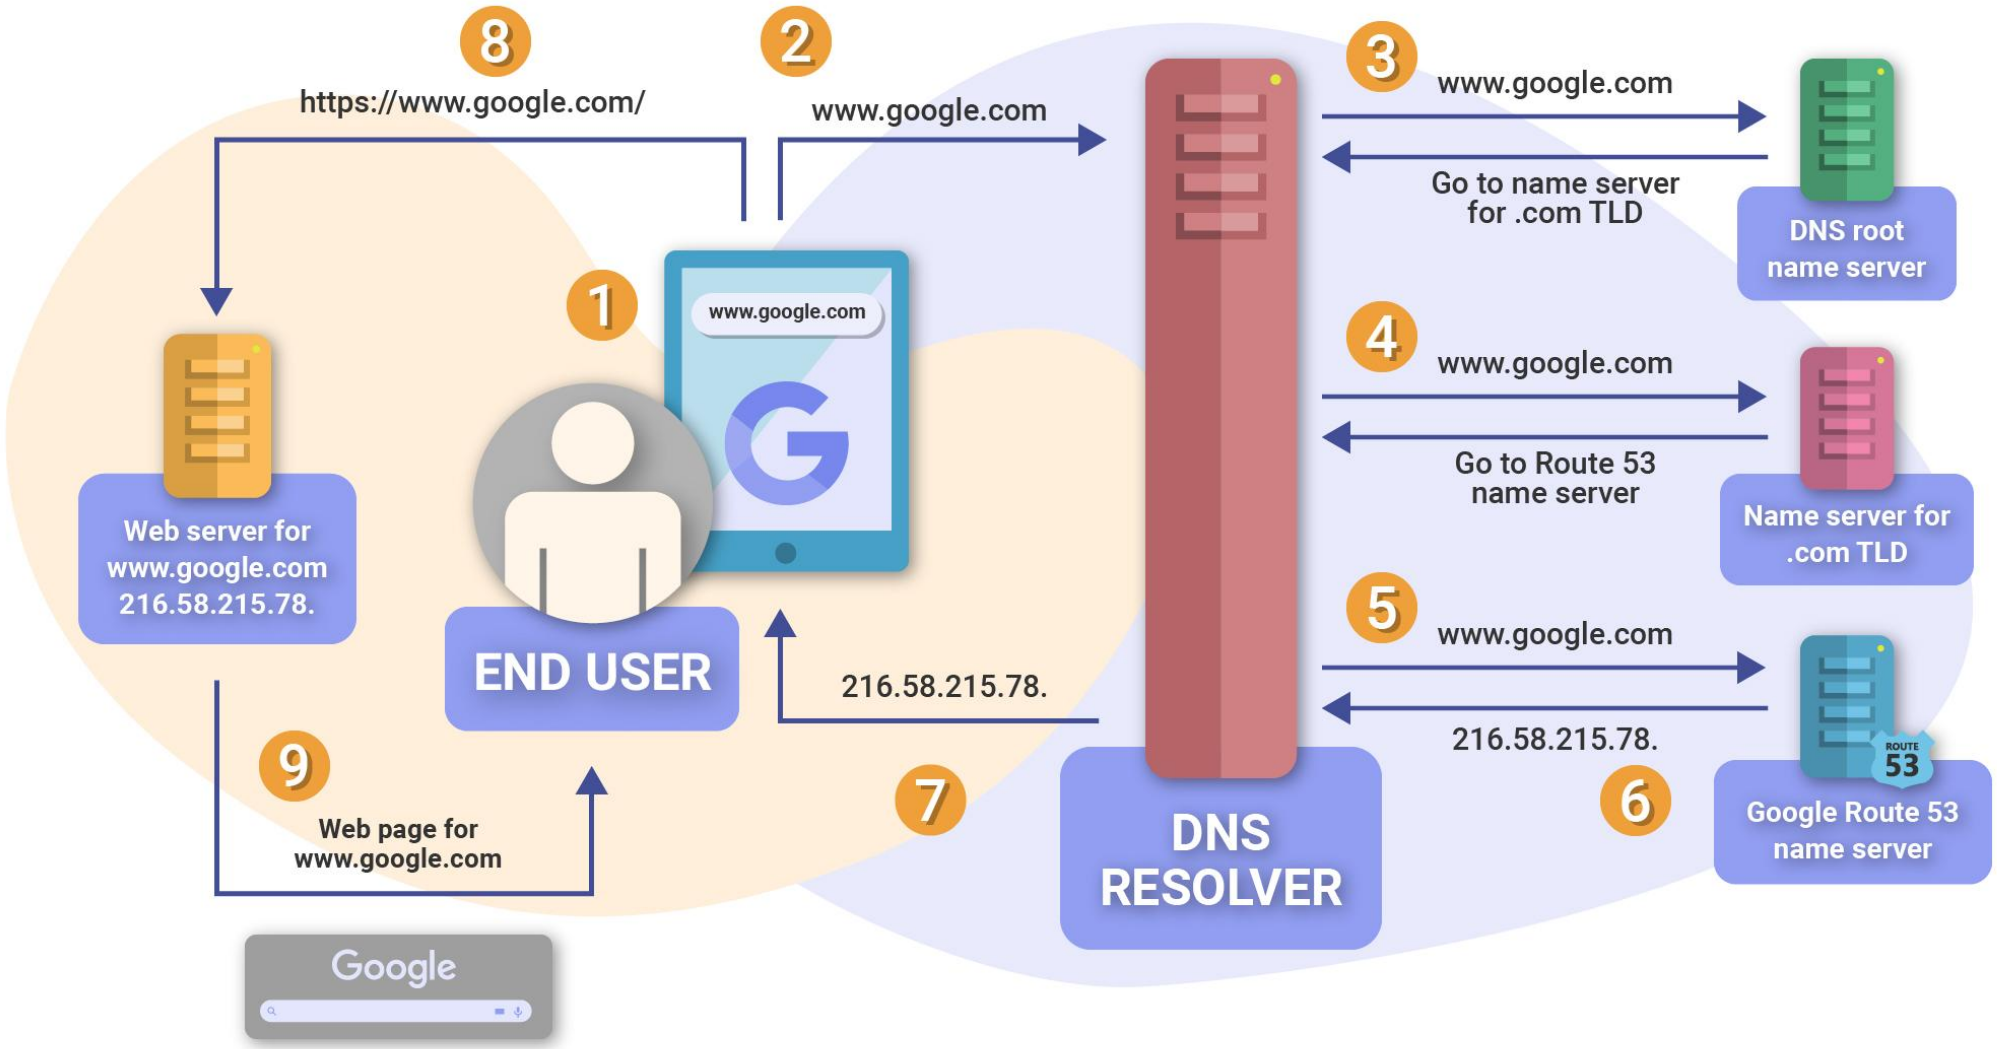

https://susodigital.com/textbook/module/wonders-of-the-world-wide-web/how-the-internet-works

TLD and URL: Quick Glance

URL: Uniform Resource Locator, we know it as a website address

TLD: Top-Level Domain like `.com` or `.edu`

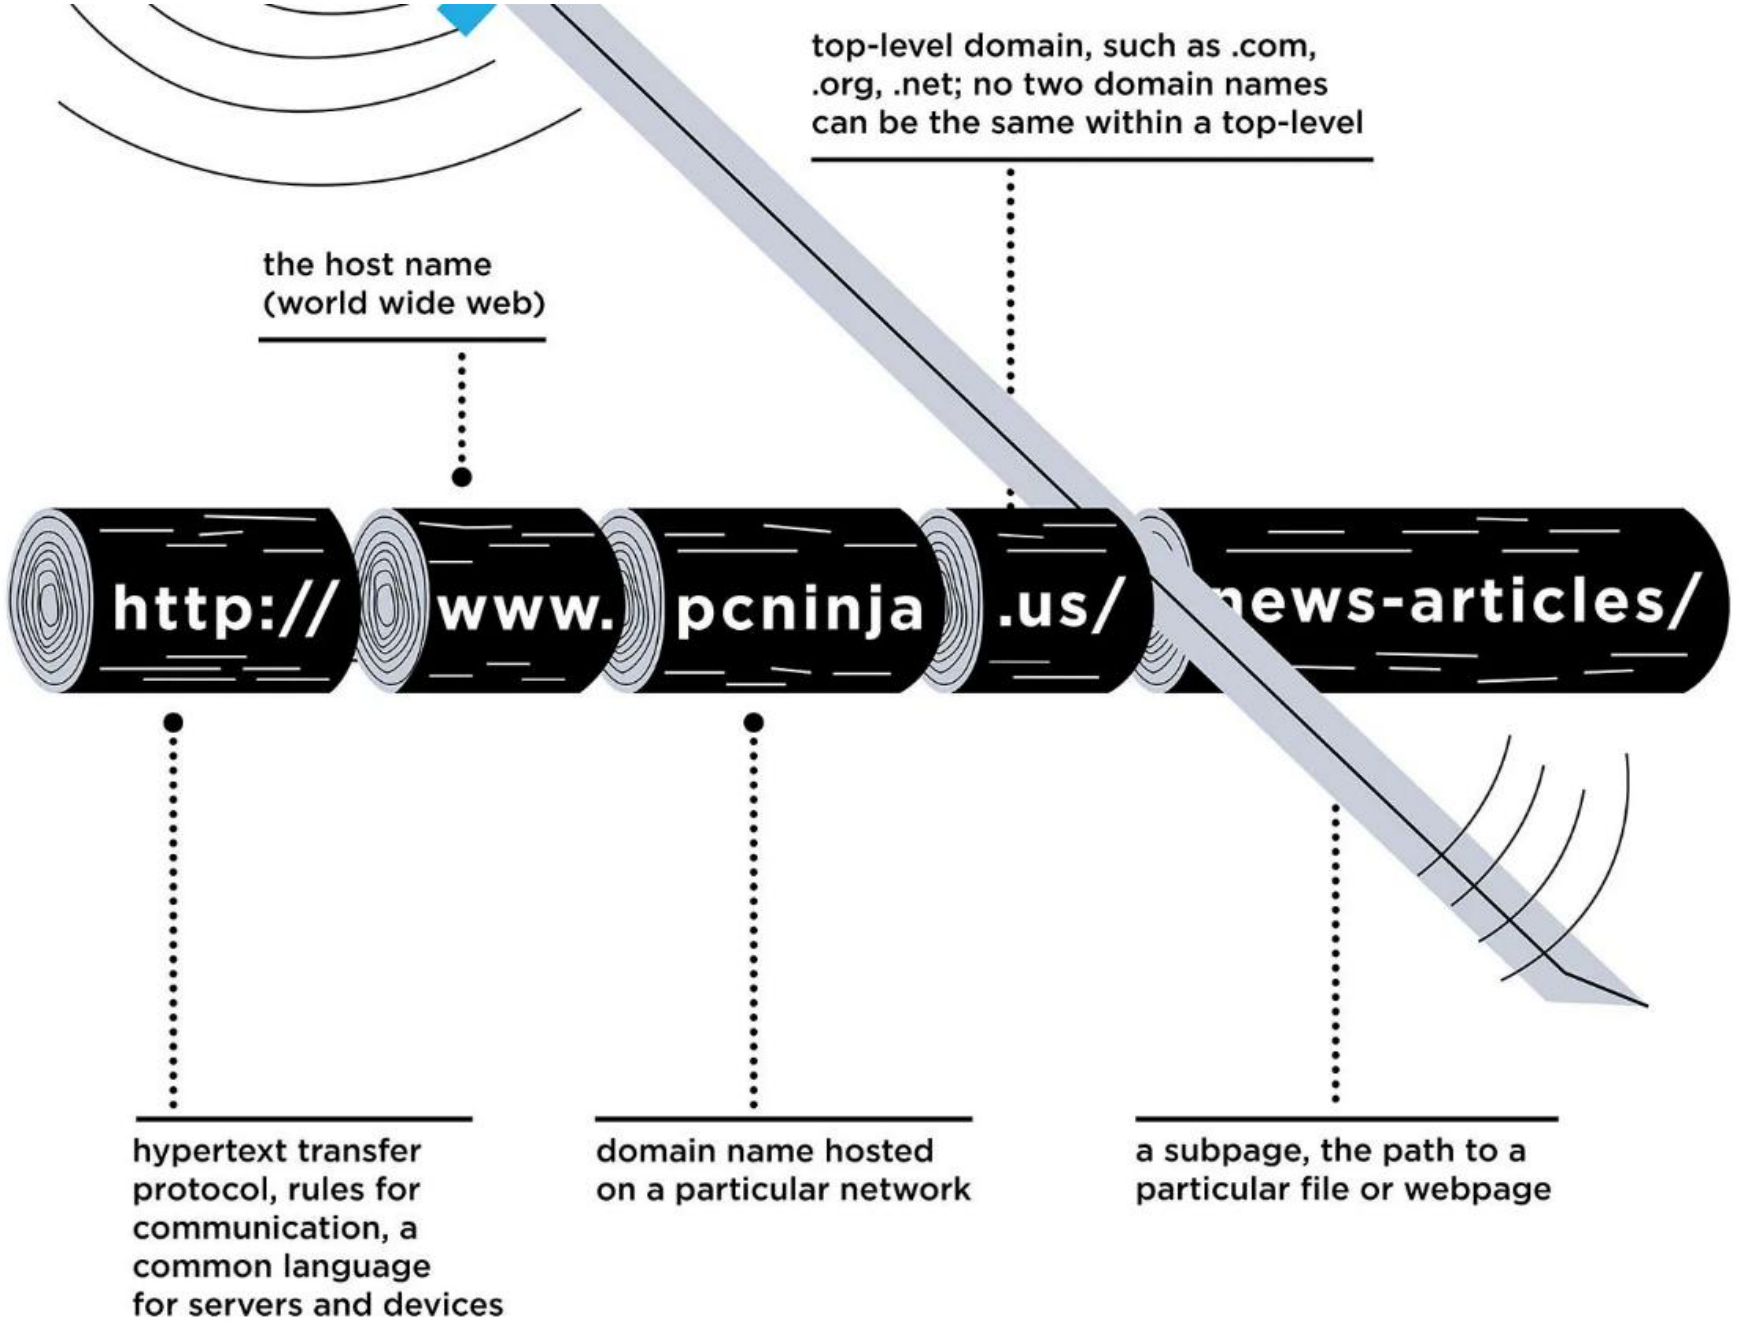

https://helloitsliam.com/2014/12/20/how-the-internet-works-infographic/

HTTPS and Port Number

- After your browser gets the `chatgpt.com` IP address, it needs to use **HTTPS** (Hypertext Transfer Protocol Secure) and a combination of IP and *port number* (a number between `0` and `65,535`) to begin searching for the actual resources that belong to `chatgpt.com` webpage
- Think of a port number as a *door* to a specific program on a device, and the program communicates on the network through that door; each device can have up to `2^16` port numbers (`2^16 = 65,536`)
  - Does it mean you can open `65,536` tabs?! Nope, browsers limit it to only several thousand...
- Each **browser tab** in this case is a different **program**, and each application that talks over the Internet is a different program that needs to run on a port that has not been taken by another program
  - Your operating system (Windows, Linux, MacOS, Android, iOS) helps manage all of that

IP + Port = program's identity on the device

- Let's say your device's local IP address is `192.168.0.2`
- Let's say that the browser's tab with `chatgpt.com` got a port number `2025` assigned by the operating system
- Then your browser will send a request to your home router, saying that it wants to find `172.64.155.209:443` and that if it has been found, then please return the result to `192.168.0.2:2025`
  - Note that the port number is now attached to the IP address via the colon
- Now your program (a specific browser tab) has an *identity* (IP + port), and your home router will map that identity to its public (external) IP and a random port, like `23.0.11.57:11111`
- From now on, your original packet will be known as `23.0.11.57:11111` during the rest of its journey across the Internet in search of `172.64.155.209:443`
- This is where *routing* algorithms kick in to determine the optimal path between your device and `chatgpt.com` server that *serves* (hence, it's called a *server*) the website files that are later shown in your browser

*Side Note: Reserved Ports

- As a side note, there are many ports that are “reserved” to be used by certain programs that do a common task on the network
- For instance, if a program resolves DNS queries, it will be *listening* on port `53` (for secure DNS, `853`)
  - **Listening** means the program is visible (discoverable) on the network, and others can send packets to it and receive a response
- Another example: if a program serves websites over the **HTTPS** protocol, then it will be listening for connections on port **`443`**
- Reserved ports do not necessarily mean they cannot be used for something else, but it *does represent a standard* that many application developers follow

Routing (From Home to ChatGPT)

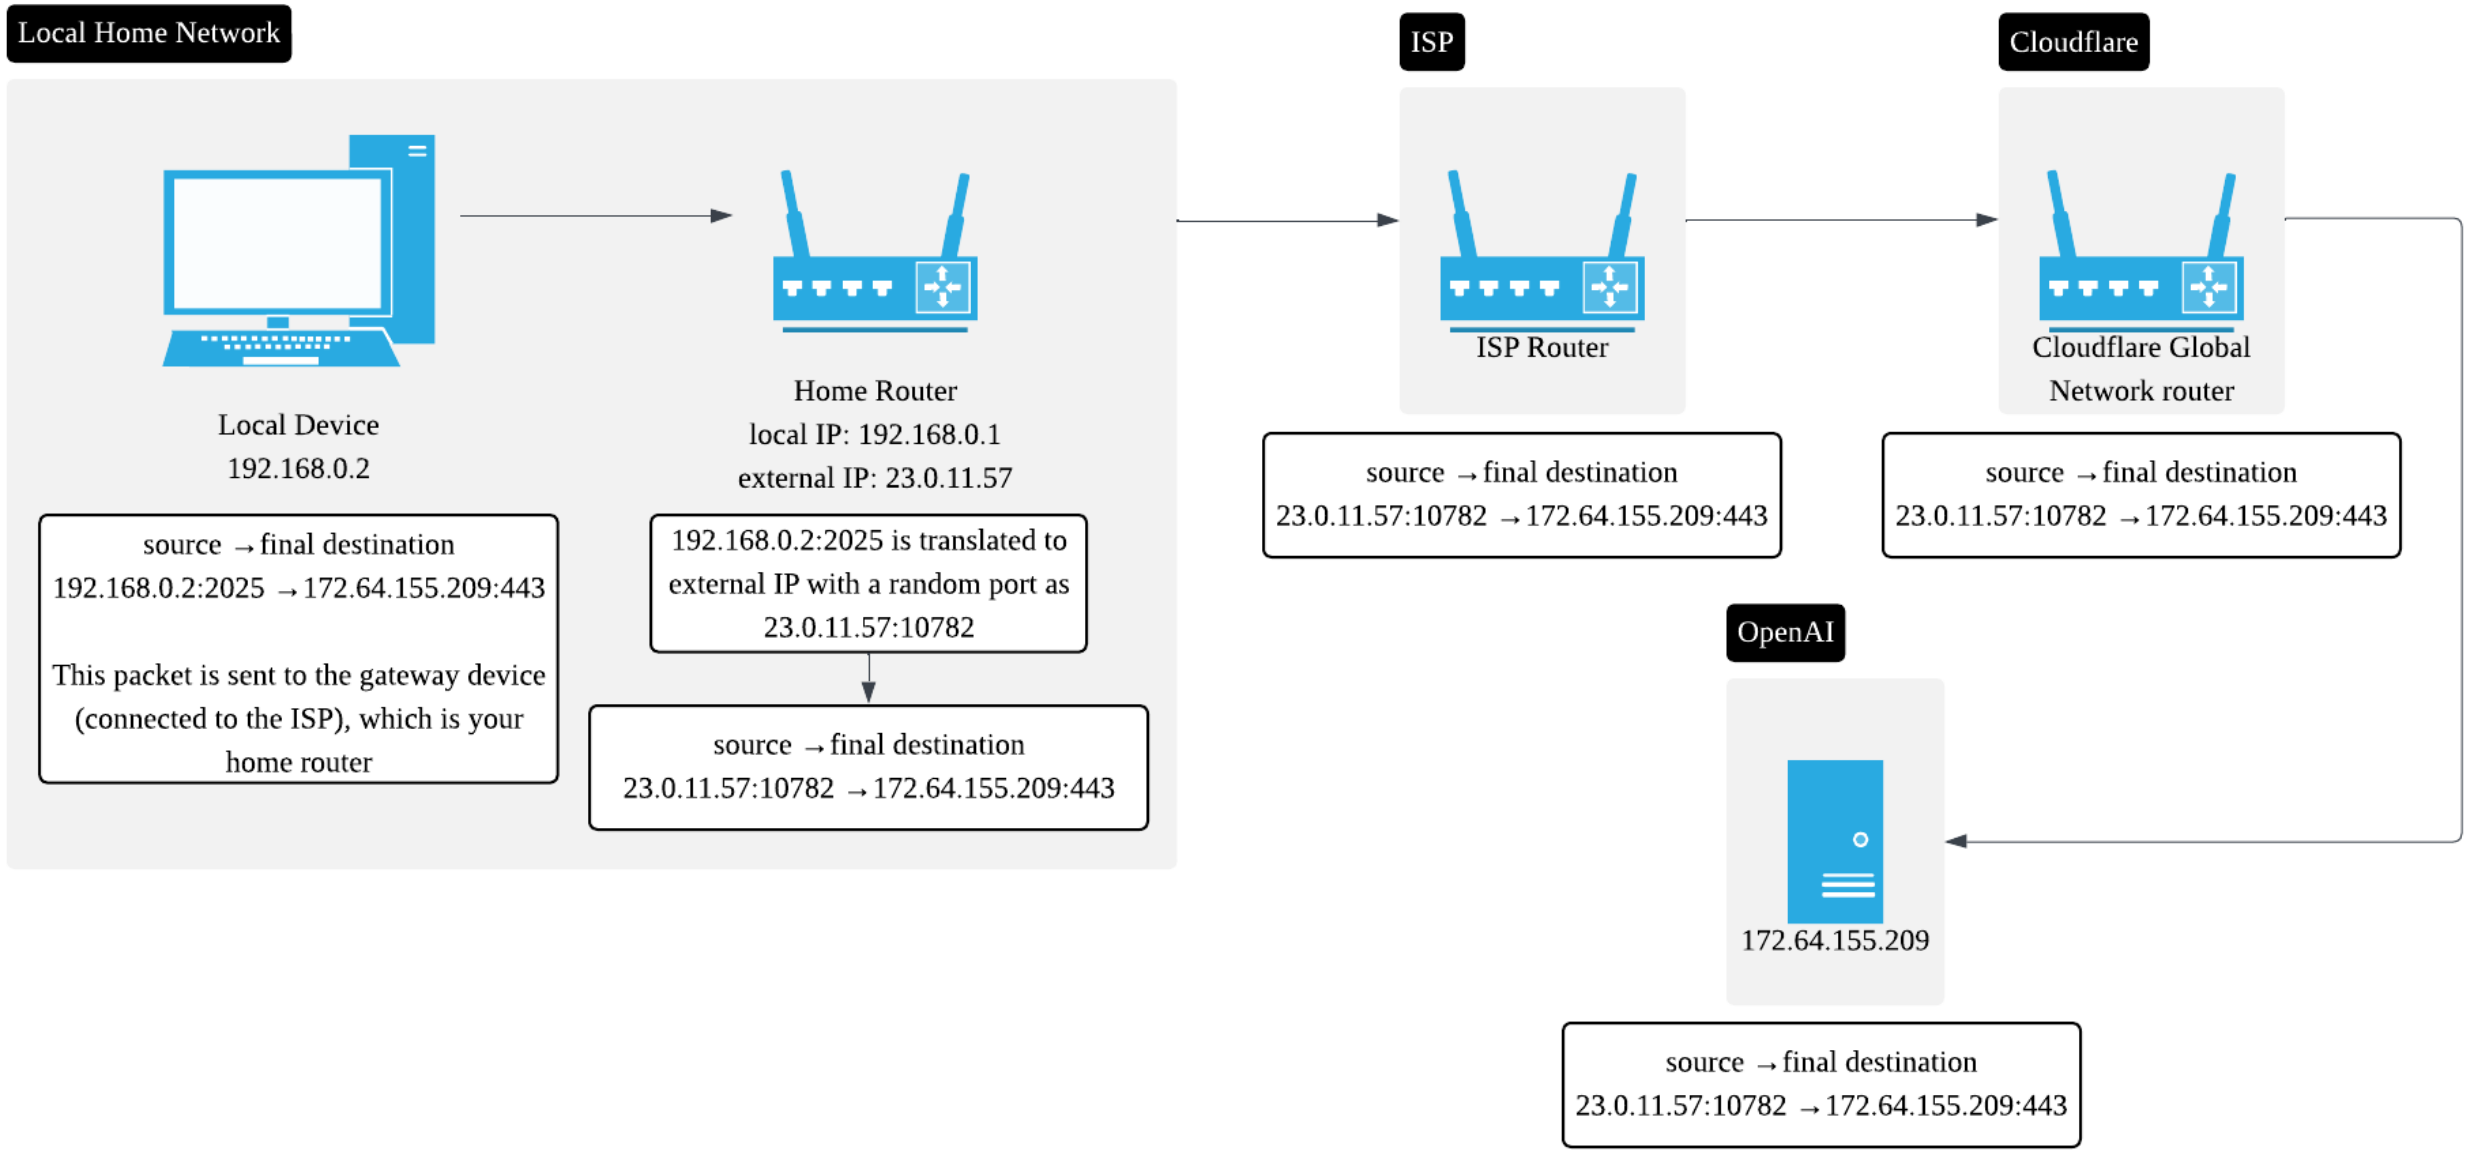

Routing

- Every time your original request travels to the next routing device (aka *hop*) on the Internet, that device looks up the packet's destination and sends it to the *hop* that it believes is the closest to that destination
- On every *hop,* starting from the ISP, the routing device (aka *router*) knows where the packet is going to (`172.64.155.209:443`) and where it came from (`23.0.11.57:11111`)
- On average, it may take between `12-18` hops before your request reaches your destination
- When `chatgpt.com` server responds with a web page, the data gets split into packets and they are all sent to `23.0.11.57:11111`
- Next, your home router picks them up, looks at what they were mapped to (`192.168.0.2:2025`), and sends them to your device, where your operating system finds the program (browser tab) and forwards the packets to that program

Packet

- Your original request is often heavy, especially if your browser loads media data
- Thus, each request is split into smaller parts called *packets*, and all packets are assembled in the right order at their destination, but they can all take **different routing paths**!
- Some packets may get lost or dropped due to hardware, software, and other random issues

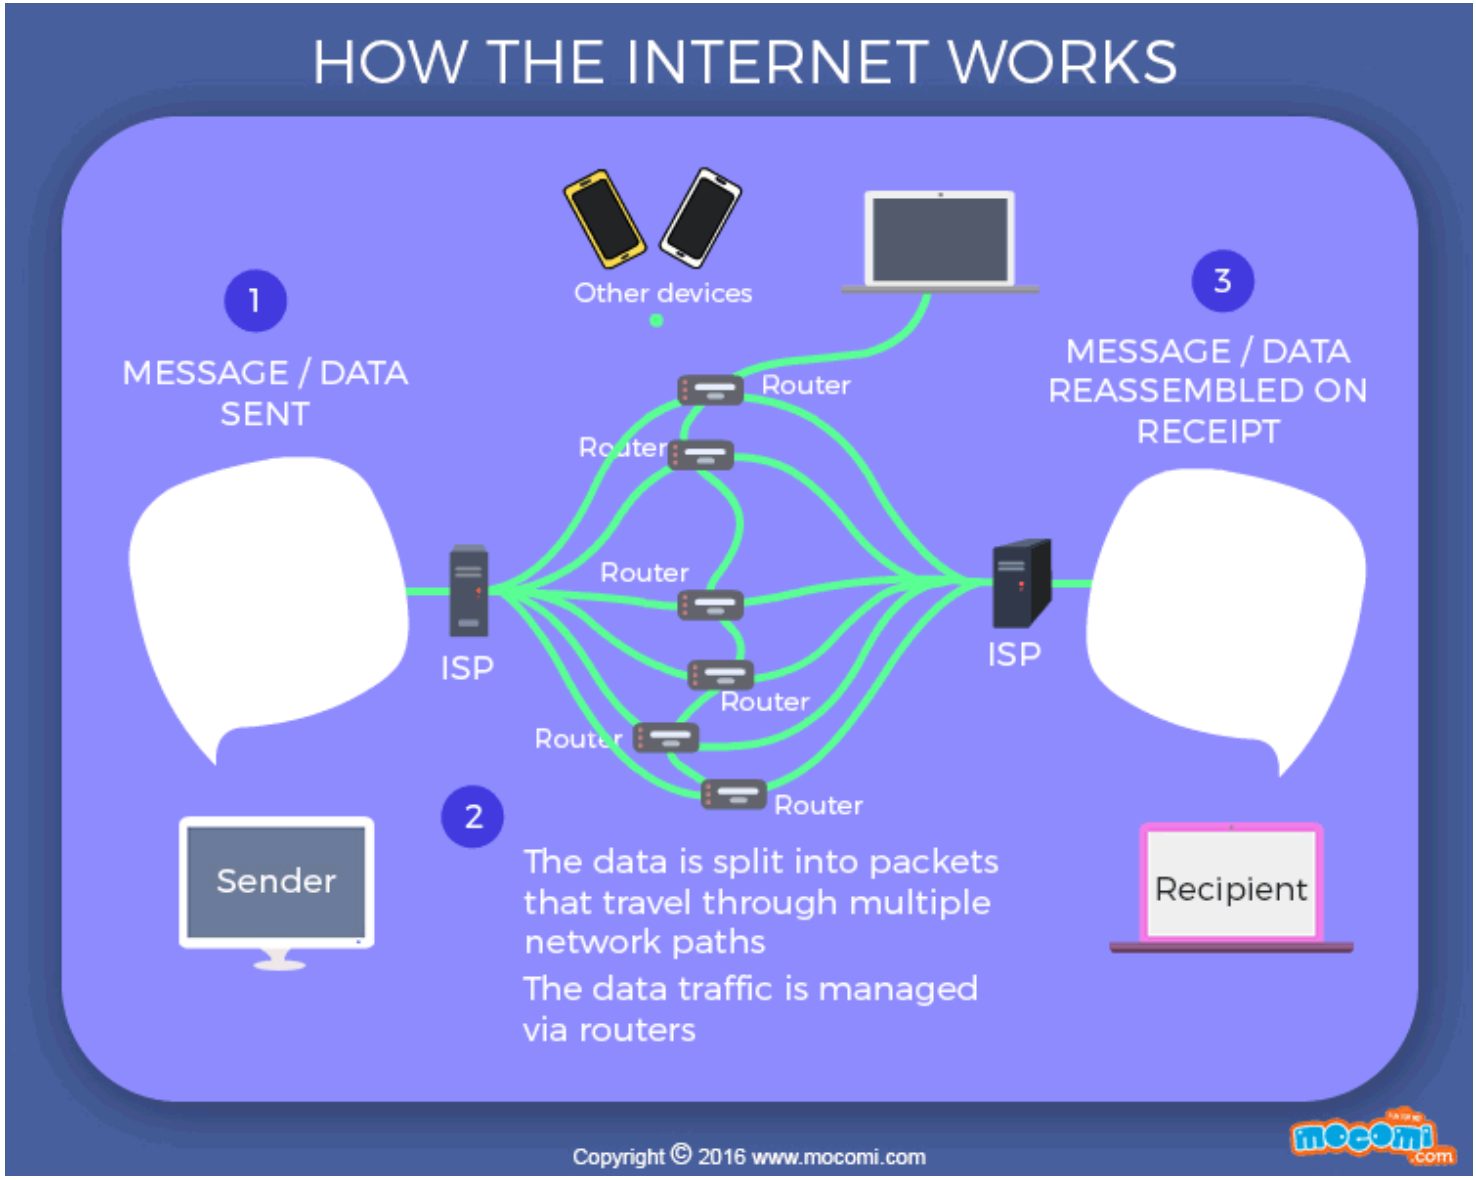

Network Problems: There Are Many

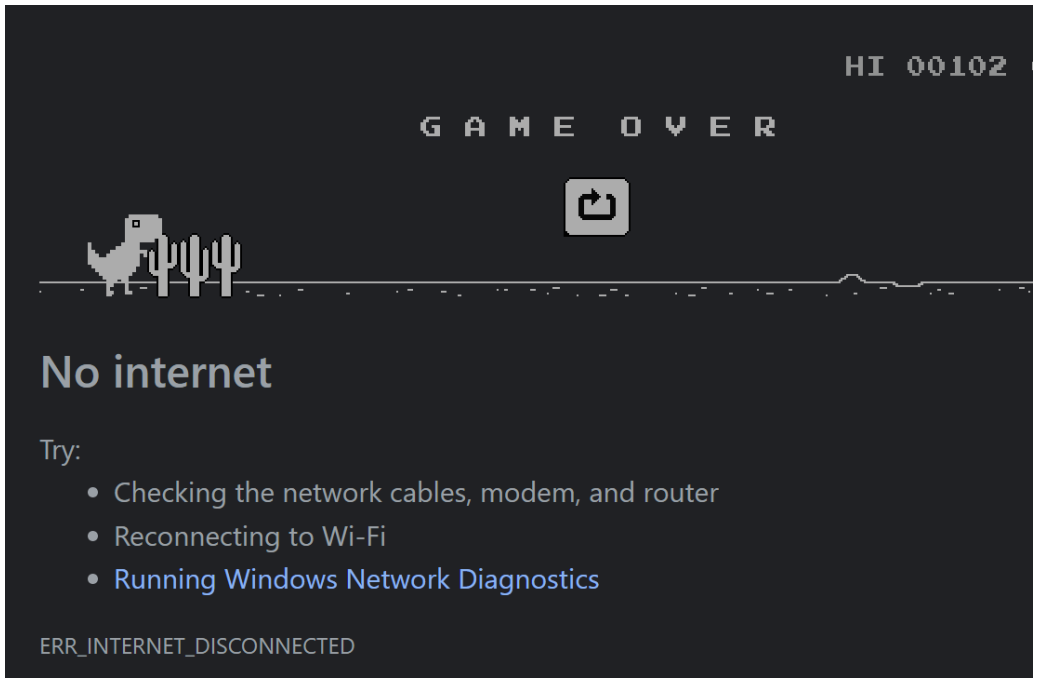

- What happens if the destination is unreachable?
  - Each packet has an attribute called `TTL` (Time To Live, what a name!!), and depending on your operating system, it can be set to `64`, `128`, or `255`
  - At each hop, the `TTL` gets reduced by 1
  - Every routing device checks if `TTL = 0`; if it is, then it deletes that packet as if it never existed
  - This mechanism prevents packets from looping around the world forever, and a simple `TTL` attribute **saves the Internet from being instantly overloaded**

- What happens if the packet is lost?
  - Depending on the protocol, your device will **try resending** it a few times and then eventually give up if no response comes back, **or it may just not even care about it** (like the *User Datagram Protocol*, UDP, used in **gaming and streaming**)

HOW THE INTERNET WORKS

SO HOW EXACTLY DOES THE INTERNET WORK? What's going on in that cloudy, tangled web? A lot of little ninja are working at super speeds to bring you the data you seek!

STEP 1 DEPARTURE

You type 'pcninja.us' into the **web browser** of your computer. Hop in the ninja mobile, and prepare for an adventure!

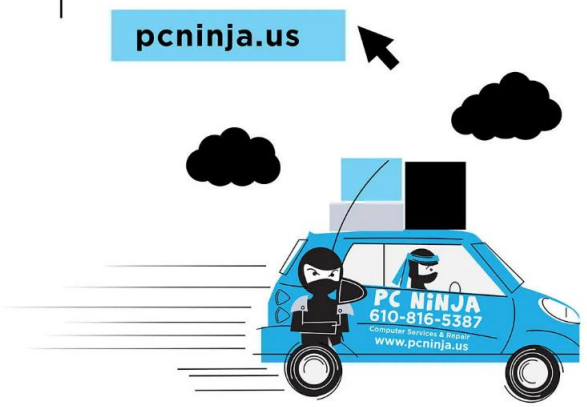

STEP 2 THE AIRPORT

Your computer is connected to the Internet through a **modem** and/or **router**, a jumping-off point to other networks in the world. **Firewalls**, in your browser and/or modem, monitor incoming and outgoing data, allowing or disallowing unscrupulous data on the network. Airport security ensures safe travel.

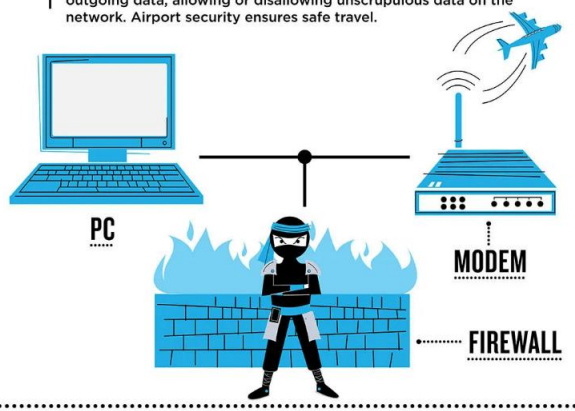

STEP 3 CUSTOMS

Your query is submitted to your **ISP (Internet Service Provider)**. Within your ISP are multiple servers to store and transfer data, a **NAP Server (Network Access Protection)**, and a **DNS (Domain Name Server)**. The DNS translates the text domain name (pcninja.us) into the numerical **IP address**. DNS uncovers the identity of the alias!

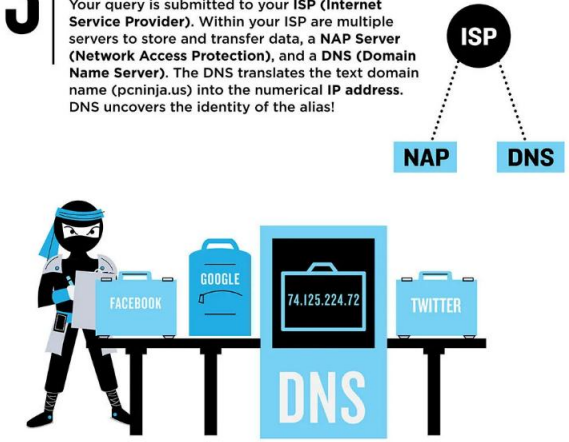

STEP 4 BUCKLE UP

Your browser maps itself to the desired IP address and establishes the hypertext transfer protocol (http), or language used to communicate on the **World Wide Web**. Good idea to take a pocket translator!

STEP 5 TAKE OFF

With the destination identified and language acquired, the ninja finds the speediest way possible to reach the server with that IP address. It travels through the **cloud**, which is made up of all the other ISPs and servers to which all the other home and business networks connect, to form that network of networks, the Internet. Not a direct flight, but no long layovers with these multiple carriers!

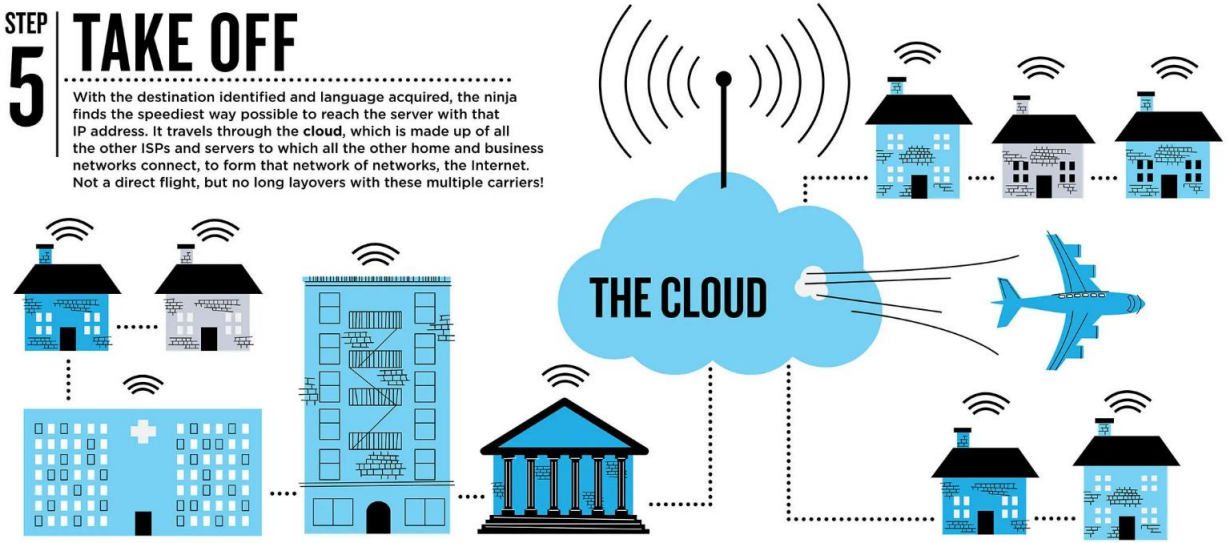

STEP 6 LANDING

Jumping from server to server on the Web, the ninja finally locates the target server hosting the target IP address for pcninja.us. A connection is established with that website and your computer. Please make sure your seats are in the upright position, we're ready to land!

STEP 7 INCOMING!

The ninja makes an even quicker return journey, bringing to your computer screen the graphical **website** of pcninja.us, which is full of data, pictures, and contact information. Thank you for flying with PC Ninja, and enjoy your browse!

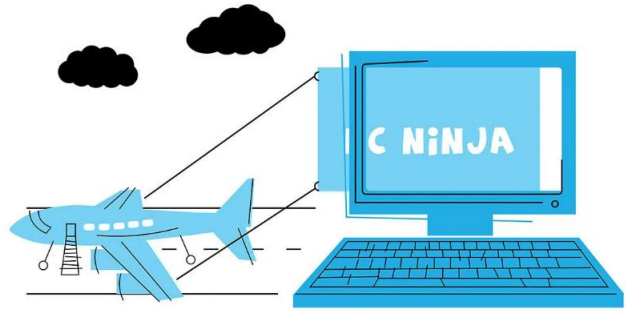

https://helloitsliam.com/2014/12/20/how-the-internet-works-infographic/